


Assignment: Paging (2 marks)

$$vaddr = \text{0x80100000}, \text{paddr} = \text{0x00100000}$$

$$vaddr = 8 \times 16^7 + 16^5 = 2^{20} + 2^{31} = 1(0)_{10}1(0)_{20} \text{ in binary.}$$

$\{ (0)_{10} = \text{ten zeroes i.e. } 0000000000 \}$



$$\Rightarrow \text{Page Directory offset} = \text{first 10 bits of } vaddr = 1(0)_9$$

$$\text{and Page Table Index} = \text{next 10 bits of } vaddr = 01(0)_8$$

~~and~~ Page Directory offset = $1(0)_9 = 2^9 = 512^{\text{th}}$ entry.
and this will contain the physical address of page table. (PPN + flags)

$$\text{Page Table offset} = 01(0)_8 = 2^8 = 256^{\text{th}} \text{ entry in page table}$$

this will contain the physical address of page ~~more~~ ~~table~~ ~~PPN~~ which is equal to paddr here

$$\text{Page offset} = (0)_{12} = 0$$

permissions \rightarrow read only, Present,

\Rightarrow

$$\text{0x00100001}$$

\hookrightarrow present flag

Assignment: Page Table Reload

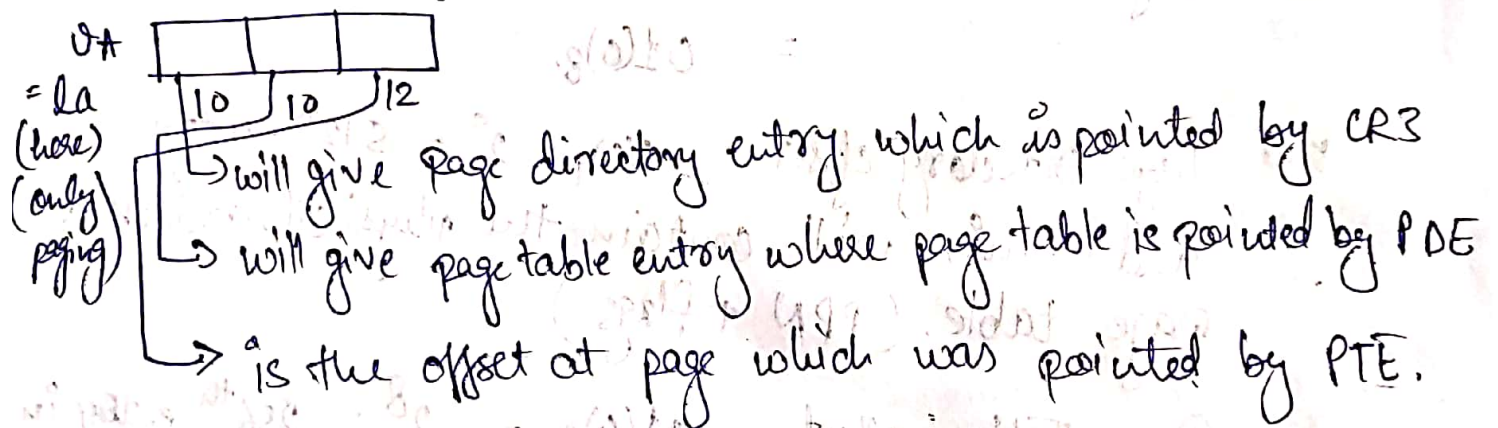
Why is kpgdir[0] zero?

Ans: $kpgdir = \text{setuptvm}()$ \rightarrow this funcⁿ only maps the space above $\text{KERNBASE}(2GB)$ not the user space.

That is why, not only 0th index but all the entries from 0 - 511 (first 512 entries \rightarrow till 2GB) were zero and from 512 index onwards there are mapped.

How would you translate 0x80107beb to a physical address?

Ans: we'll use page table and page directory to get the suitable physical address.



(gdb) `print/x kpgdir[0x200]`

\$6 = 0x3fe007

Q What is this? (PPN + flags)

Ans: Physical address of page table plus some flags. (PDE)

Q What is PPN?

first 20 bits of PDE i.e. $0x3fe007 \gg 12$

= 0x3fe

Q What does 7 mean?

Ans $\rightarrow (7)_{16} \rightarrow (0111)_2$
uses flag writable flag Present flag

(gdb) print/x ((int*)0x3fe000)[0x106]

\$12 = 0x106001

Q What is this?

Ans \rightarrow ~~Physical address of the page + flags~~
(PPN + flags) \rightarrow of page to which offset will be added.

Q Why 1 in the low bits?

1 is the present flag, showing that the page is available.

Q Why did the physical address work in gdb?

B/c we are still in entrypgdir and did not load kpgdir in cr3 yet.

Q Why it won't work after calling switchvm1?

switchvm1() will load the kpgdir into cr3 or switch to kpgdir from entrypgdir and

kpgdir does not have mapping to its own address.

Assignment : Addressing.

Ans: By adding 0x100000 to vaddr of each elf section, we are only changing the elf sections. But as functions and variables are named by addresses, the already compiled code may have internal reference to pointers having predefined locations and shifting addresses may point to wrong locations which will cause errors.

Assignment : Traps.

Q Situation in which a suspended process will have three set of saved registers in its kstack.

Ans If a timer interrupt occurred in the middle of kernel execution, it'll have 2 trapframes (for kernel & user) and 1 context set of registers.

Q 2 context and 1 trapframe

No, not possible. Context structure ~~shows~~ is used when we switch and given 2 context, where to start execution from. NO, not possible.

Q 2 trapframes and 1 context

Yes, in the situation described above where the kernel interrupt occurs in the middle of kernel execution.

Q More than 3 set of saved registers
No, to prevent kstack overflow, we disable interrupts while executing a handler ensuring atmost 2 trapframes in kernel and therefore, (2 trapframe, 1 context), we can't have more than 3 sets of saved registers.

Assignment: Context Switching

Q Where is the stack that sched() executes on?

→ sched() executes on process's kstack which is located from kernel heap space.

Q Where is the stack that scheduler() executes on?
Scheduler() executes on its own kstack, or say CPU's kstack which is initialized during kernel initialization and this does not belong to any process, it doesn't have a pgdir associated to it.

Q When sched() calls switch(), does that call to switch() ever return? when?

Sched() executes on process's kstack and when it calls switch(), that will switch to scheduler's kstack (say P1).

Now, whenever scheduler (after some process switches) picks that process again, it will return.

(and switches to P1)

Q What is the four character pattern?

→ acbad cbad cbad cbad ...

Q Why ac as very first characters?

When the kernel initializes itself, `main()` function calls the `schedular()` and it enters `schedular()` at first, therefore, a in the very beginning.

Now after a,

1. `cprintf('a');`

we switch → 2. `switch(&cpu → schedular, &proc → context);`
to process `proc`. 3. `cprintf("b");`

Then some process calls `sched()`, therefore, we see "c" as second character on screen.

now we switch 4. `cprintf("c");`

to `schedular` → 5. `switch(&cpu → schedular, &proc → context, &cpu → schedular);`
again. 6. `cprintf("d");`

↳ In `schedular`, we'll continue from where we left, i.e. line 3: `cprintf("b")`, and as there's a infinite for loop in `schedular()`, we again come to line 1: `cprintf("a")`, switch from here, now whenever `sched()` is called, we continue from line 6, i.e. `cprintf("d")` and this continues
⇒ pattern is a cbad cbad cbad ----- so on.

Q Could switch do less work and still be correct?
Can we reduce the size of struct context?

In general, as per the gcc calling conventions, switch requires to save the callee saved registers and doing less work, i.e. not saving these registers can only work in case the register values are not changed, which can't be true for all the function calls.

We can definitely make the context to point to the last register value i.e. edi and make struct contain only this value but save the rest of registers value in backup. This way, we reduced context struct size but followed gcc calling conventions too.