

⇒ Assignment: Paging

va: 0x80100000 → pa: 0x00100000  
↓

top 20 bits: 1000 0000 00 = 512

next 10 bits: 0100 0000 00 = 256

512<sup>th</sup> entry of page directory will point to a page table,  
256<sup>th</sup> entry of that page table will point to physical  
address 0x00100000. Permissions: PTE-P (PTE-W is  
not set to make it read only)

⇒ Assignment: Pagetable reload

(gdb) print/x kpgdir[0]

→ why is this zero?

↳ At this point, kpgdir is setup for the kernel  
part of the page table by setupkvm() function.  
It has mapped pages above KERNBASE address,  
and 0 to KERNBASE is still empty and available  
for user processes. That's why kpgdir[0] is zero.

→ how would we translate 0x80107beb to a physical  
address?

↳ (for x86) V2P-WD() function can translate <sup>virtual</sup> ~~physical~~  
address to physical address. We can do it  
by subtracting 0x80000000 (KERNBASE).

$$0x80107beb - 0x80000000$$

$$= 0x107beb$$

```
(gdb) print/x 0x80107beb >> 22
```

```
$4 = 0x200
```

```
(gdb) print/x kpgdir[0x200]
```

```
$6 = 0x114007
```

→ What is this ?

↳ 2nd level page table address.

→ What is PPN ?

↳ 0x200

→ What does the 7 mean ?

↳ Last three bits are '1'. They corresponds to flags present, writable and user. Here, page is present, writable and can be accessed by all.

```
(gdb) print/x (0x80107beb >> 12) & 0xfff
```

```
$6 = 0x107
```

```
(gdb) print/x ((int*) 0x114000)[0x107]
```

```
$12 = 0x107001
```

→ What is this ?

↳ This is starting address of the page containing memory location we want to access.

→ Why 1 in the low bits ?

↳ It indicates that the page is present.



→ why did physical address work in gdb?

↳ because paging is not enabled yet. (flag in cr3 register.) ~~There is~~

After switchkm loads kpgdir into %cr3

(gdb) x/i 0x107beb

0x107beb: can not access memory at address 0x107beb

→ why?

↳ switchkm enables paging by setting the flag in cr3 register. Now, all addresses need core virtual and can be ~~accessed~~ accessed via paging only.

## ⇒ Assignment: Addressing

Suppose you wanted bootmain() to load the kernel at 0x80200000 instead of 0x80100000, and you did it by modifying bootmain() to add 0x100000 to va/pa of each ELF section. Something would go wrong. What?

↳ bootmain() calls entryfunction which maps 0 to 4MB and 2GB to 2GB+4MB of va space to 0 to 4MB of pa space. After this pa > 4MB can't be accessed. If kernel is loaded at 2GB+2MB instead of 2GB+1M there might be loss of information.

## ⇒ Assignment : Traps

→ Is it possible to have two "context" structures and one "trapframe" structure on the kstack? If so, when? If not, why not?

↳ NO. kstack can not have two "context" structures. Because, after pushing context on a kstack, you either pop the same context or ~~switch~~ pop context from some other process' kstack. ~~Therefore~~, ~~there~~ Interrupts are also disabled during this operation. ~~Therefore~~

→ Is it possible to have two trapframe structures and one context structure on the stack. If so, when? If not, why not?

↳ YES; When there is an interrupt in user process, it starts running interrupt handler in kernel mode after storing <sup>on kstack.</sup> trapframe<sub>x</sub>. While running in kernel mode, if a timer interrupt occurs then kernel values will be stored and switch to another process. In this case, there are two trapframes (user trapframe and kernel trapframe) and a context structure on the kstack.



→ Is it possible to have more than three sets of saved registers in the stack? If so, when? If not, why not?

↳ NO. for xv6, kstack can contain maximum two trapframes only. It is ensured by using locks and disabling external interrupts, which such that kernel instructions are not interrupted in the middle.

## ⇒ Assignment: Context Switching

→ where is the stack that sched() executes on?

↳ sched() calls switch, it executes on kstack of currently running process.

→ where is the stack that scheduler executes on?

↳ scheduler doesn't have its own address space or stack. It runs on the same stack on which main() function is running.

→ when sched() calls switch(), does that call to switch() ever return? If so, when?

↳ Yes, switch() will return, but not immediately.

It will return when the process is resumed next time.

→ Could switch do less work and still be correct?  
Could we reduce the size of a struct context?  
Provide concrete examples if yes, or argue why not.  
↳ NO. To ensure correctness, callee-save registers need to be stored to maintain gcc calling conventions. Therefore, size of struct context can't be reduced and switch can't do less work.

→ What is the four-character pattern?

↳ Output: acbadcbadcbad...

repeating pattern: "cbad"

→ The very first characters are ac. Why does this happen?

↳ While setting up processor, scheduler is initialized for the first time. therefore 'a' is first character.