

Question 1 HW-3
Paging

SALIL CHANDRA
2016 EE30246

Ans. For page directory:

① Offset = $0 \times 200 = 512^{th}$ entry

② Entry value = $0 \times \underbrace{00200,005}_{\substack{\text{Can be any} \\ \text{Physical address of any} \\ \text{allocated page}}} \rightarrow 0 \ 1 \ 0 \ 1$

↓ Present
↓ Read only
User

For page table (2nd level):

① Offset = $0 \times 100 = 256^{th}$ entry

② Entry value = $0 \times \underbrace{001000005}_{\substack{\text{PTE-U} \\ \text{PTE-W} \\ \text{PTE-P}}} \rightarrow 0 \ 1 \ 0 \ 1$

↓ Write back
↓ User
↓ Present
↓ Read only

Question 2. Page Table Reload

① (gdb) print/x kpgdir[0]. Why is this zero?

Ans. All entries in kpgdir are initialised to zero in setupkern() (line 1825: memset(kpgdir, 0, PE_SIZE);). Then mappings are made only starting from KERNBASE. Therefore ~~entries~~ entries below KERNBASE are still 0 (unmapped) and are to be used for user address space mapping later.

② How would we translate 0x80107beb to a physical address?

Ans. ~~First 10 bits are used to offset into PD, next 10 bits offset into PT. Final 12 bits are offset into the physical mapped page in the physical address space.~~
Since setupkern not called yet, $PA = VA - \text{Kernbase}$, $0x80000000$

③ (gdb) print/x 0x80107beb >> 22

\$4 = 0x200

(gdb) print/x kpgdir[0x200]

\$6 = 0x114007

(i) What is this?

Ans. $\$4 = 0x200$ are the first 10 bits of the virtual address that index into the page directory.

$\$6 = 0x114007$ is the page directory entry at offset 0x200 (ie, 512th entry).

(ii) What is the PPN?

PPN is for page table (2nd level page table page) is 0x114000.

(ii) What does 7 mean?

Ans. 7 (0111) is the last 4 bits of the PDE, which implies the page is Write back, User accessible, Writable and Present.

④ (gdb) print /x (0x801076eb >> 12) & 0xfff
\$6 = 0x107

(gdb) print /x ((int*) 0x114000) [0x107]
\$12 = 0x107001

Ans. \$6 = 0x107 is the offset into the 2nd level page table (263rd entry).

\$12 = 0x107001 is the 2nd level page table entry at offset 0x107. (0x107000 is the PPN, 001 are flags)

(ii) Why 1 in the low bits?

Ans. 1 in low bit indicates that the page is present. (0th bit is the PTE-P flag)

⑤ Why did the physical address work in gdb?

Ans. kpgdir is still not being used because switchmem() has not been called yet. Since presently, the virtual addresses from (0 to 4MB) and (Kernbase to Kernbase + 4MB) are identically mapped to physical addresses 0 to 4MB (due to entrypgdir); we can still access the physical addresses directly.

⑥. Cannot access memory at address $0x107\text{ keb}$.
why?

Ans After `setupkm()` is called, the paging hardware starts to use the `kpgdir`. In `kpgdir`, only addresses above `kernbase` are mapped. Since $0x107\text{ keb}$ is an address below `kernbase`, it cannot be accessed.

Question 3 Addressing

Ans . The ~~xx~~ addresses may overflow above 4MB. Since entrypger only maps ^{VA from} (0 to 4MB) and ($Kbase$ to $Kbase + 4MB$) to (0 to 4MB) in the physical address space, ~~if all addresses~~ (assuming XVB fits inside ~~the~~ 4MB space), if we offset all addresses by 1MB, they may overflow the 4MB cap, and therefore ~~overflow~~ will not be valid while using entrypger .

Question 4. Traps

- ① Explain a situation when a suspended process will have 3 sets of saved regs in its k-stack.

Ans. Suppose a process makes a system call (Therefore user registers are saved in trapframe → ①). When trap() is returning, an external ^{timer} interrupt occurs and the kernel side registers for the process get saved in another trapframe lower on the same stack. ~~Sup~~ Now the timer interrupt will cause context s/w, therefore kernel registers will get ~~to~~ saved again in struct context.

- ② Is it possible to have 2 'context' structs and one 'trapframe struct' on k-stack.

Ans No, there can't be 2 context structures on the same k-stack because the ~~same~~ interrupts are disabled while switching, therefore no timer interrupt (which causes struct context to be pushed) doesn't appear.

- ③ Is it possible to have 2 trapframes and one context structure on the same k-stack.

Ans Yes. Exactly same scenario as in ①.

i.e., the chronology of system call → timer interrupt (while returning from sys call) → context switch.

④. Is it possible to have ~~the~~ greater than 3 sets of saved regs on the k-stack?

Ans No. Since there can be a maximum of 2 trapframes at a time on the stack (system call followed by external interrupt). There cannot be a 3rd trapframe because interrupts are disabled while handling the external interrupt.

Also, a maximum of 1 context structure can appear on the k-stack at a time (as in ②).

Therefore >3 sets of saved regs are not possible.

Ques 5. Context Switching

①. Where is the stack that `sched()` executes on?
Ans - `sched()` executes on the currently running process's k-stack.

②. Where is the stack that `scheduler()` executes on?
Ans. `scheduler()` executes on the main stack (where the `main()` function is executing).
However in certain instances of process model, the scheduler can have a separate stack, but not in xv6.

③. When `sched()` calls `switch()`, does that call to `switch()` ever return? If so, when?

Ans. Yes, `switch()` function returns when the process is again 'context switched in', which may not happen ~~at a later~~ immediately, but at a later context switch.

④. Could `switch()` do less work & still be correct? Could we reduce size of struct context? Provide examples.

Ans. No, ~~size of~~ `switch` cannot do ~~less~~ less work and we cannot reduce the size of struct context.
This is because `switch()` only saves the callee-saved registers (according to gcc calling convention), which is the bare minimum to ensure correctness.
~~Struct context must save all registers~~

⑤. What is the four char pattern?

Ans. cbad

(O/P is "a cbad cbad...")

⑥. The very first characters are 'ac'. Why?

Ans.