

HW-3

PARTY CHOPRA

2016ME10323

OPERATING SYSTEM

411 x86 4KB pages (Paging)

VA
0x80100000

PA
0x00100000

Translation Overview

DIR PAGE OFFSET

PAGE DIRECTORY

PAGETABLE

DIR ENTRY

PG TBLENT

PAGE FRAME

PHYSICAL ADDRESS

CR3

0x80100000

0x200

PAGE DIR OFFSET

0x200

PAGETABLE
OFFSET

0x000

Physical Page
offset

ENTRY at in PAGE DIRECTORY

$(CR3 + 0x200 * \text{size of PTE})$

4 bytes

MODIFIED TO

not specified

11 12 FLAG
BITS

0x00384

0 0 1

Read Hint Bit
Reserved Bit

Entry is at PAGETABLE 0x00384

MODIFIED TO

$0x00384 + 0x200 * 4 \text{ bytes}$ to

0x00100 001

A₂ PAGE TABLE RELOAD.

(i) Virtual Address starts from 0 to $2^3 - 1$ (Index of 32 bit machine) when we addressing 1st instruction we access it by `kpdin[0]` This gives returns 00

(ii) Just like question to translate 0x80107beb to PA

10	10	12
----	----	----

- We use first 10 bits for finding offset from CR3 pointer
- Next 10 bits used to find offset from page table base to find pointer to 4kb aligned page.
- Last 12 bits find exact location of addressed memory in the page, i.e. page offset.

This we can translate from P VA to PA once paging is enabled and pages being 4kb.

(iii) 0x114007.

* This is page table directory entry which is used to get the address of ~~the~~ page table. where ~~0x80107beb~~ ^{table base}

(iv) 0x114 is PPN.

(v) 7 in the end corresponds to the flags for the page table as a whole.

7 \equiv 1 1 1

→ presence bit

→ Read/Write bit

→ User accessible bit

However all this is provided, The specific page also has all these bits set.

Here all three bits are set.

which means page table present. it readable/writable.

User can access the page

The specific page also

(vi) This is page table entry which is used to get to the page containing the data associated with VA. First 20 bits tell us about the address of page. as all pages are 4kb aligned (4kb in size so in continuous memory, so 4kb apart) last 12 bits used as flags associated with the page.

(vii) Here 1 is set in lowest bits indicates the flags look like 001 which indicates that this page is present. moreover is read only mode. It takes kernel priv to access the page. user can't access the page.

(viii) Physical Address worked in x86. because x86 uses that structure to address the PA and VA hence. PA also worked here and so we can use here.

(ix) However, once kernel calls `setupkern` to create page table to start using it. `setupkern` adds page `0x107` into `0x3`.

Now `0x107` can't be accessed in user mode. Cause the flag is set such that only kernel will be able to access it as we are using paging to ~~access~~ ^{reach} the address.

A₃

Addressing

This will be wrong because once the kernel executing starts it will set up paging hardware to map virtual address $0x80100000$ to physical address starting at $0x00100000$; kernel assumes that there is physical memory at lower address. At this stage boot process, paging is not enabled. Instead kernel to specifies that ELF starts at $0x00100000$ which will cause boot loader to write what makes sense when paging hardware will eventually point to these crashing the system.

A₄

Traps.

- (i) It is not possible for a suspended process k-stack to have two "context" structure because once a system is context switched out only it returns to running state only when using the context structure. Moreover, a process can be context switched out once. Additionally, context structure is fast entry on the process k-stack.
- (ii) Yes, it is possible to ~~two~~ have two trap frames structure along with context structure on k-stack, this can happen if for a process the first trap occurs due to software interrupt / exception / ^{or} and second due to external interrupt received while executing in the kernel. While handling this it can be context switched out, resulting in two trap structure and one context structure.

ii)

No on a process k-stack can have σ_{max} context structure saved registers and '2' trap structure saved registers so can't have more than 3.

A5

(i) process k-stack sched() executes on.

(ii) Each CPU scheduler has its own stack.

(iii) When sched() calls switch(), this switch function does the job of switching between two contexts. The switch() returns to a different process. At a later stage again switch would be called it might return to this sched() and then this process would be resumed again caller doesn't get to see return.

(iv) Yes, switch could do less work and even then be correct. when timer interrupt occurs ^{or queue's on its own} process yield() which calls sched() which in turn calls sched switch(), assuming GCC conventions, already ~~on the k-stack~~ there are useful values of registers saved, so really it can work with that. no need to re save and pop values can just switch(esp)(eip)

(v) ac db (cabd) cabd In one core this will remain consistent like this.

(vi) This happens because 'first process' created doesn't actually context switch out, it doesn't call switch the k-stack is initial made as if the init process was ready to run without in reality ever being contexted out. Hence very first time 'i.ac' appears we start a Single CPU system.