

APRIL 2020						
Wk	M	T	W	T	F	S
14		1	2	3	4	5
15	6	7	8	9	10	11
16	13	14	15	16	17	18
17	20	21	22	23	24	25
18	27	28	29	30		

Entry: 2013CS50286
Anket Prakash Hirulkar

2020

Homework-4

Wk 11 • Day 070-296

March
Tuesday

10

Assignment: Spin-locking

09

Q: with `sti()` after acquire and `cli()` before release,
the kernel panicked.

→ There was a clear message on shell showing "panic"
after the above changes were made.

CL (Clear interrupt flag) disables any upcoming
interrupts.

01

Here there is race condition in ide.c because of
which the kernel panics.

The atomicity is not maintained

The kernel panics on second acquire().

Q: With changes in file.c [`sti()` after acquire() &
`cli()` before two releases] the kernel does not
panick.

→ In file-table-lock, the race condition does not
occur.

There is atomicity in instructions achieved.

first release cleared `lk → PS[0]` and `lk → CPUF`
before clearing lk-locked.

Since the interrupts were disabled with `cli()`,
second release ran only after first release
completed its execution.

After that second release ran into execution.

APRIL

MAY

JUNE

11

March
Wednesday

Wk 11 • Day 071-295

10	2	3	4	5
11	9	10	11	12
12	16	17	18	19
13	23	24	25	26

• Uniprocessor Locking:

① `LOCK(L) {`
 `cli(); // disable preemption`
 `while (L == 0) continue; A`

11 `L = 0;`
 `sti(); # enable preemption`

12 `unlock(L) {`

01 `L = 1;` B

02 } }

02 \Rightarrow This implementation will not work.

03 consider thread 1 was at step A where
 it had to check if $L \neq 0$ (unlocked)

04 Assuming L was already 0, thread 1 continued.

05 But at the same time, thread 2 was at
 step B so it set $L = 1$, by the time
 thread 1 moved ahead to set $L = 0$.

07 then thread 2 made an error of judgement

Due to this discrepancy (potential discrepancy),
 this lock will not work.

8

Since both threads execute $L = 0$ & $L = 1$ at the
 same time - thread 2 would unlock the lock
 nullifying effect of thread 1.

6 7 8 9 10 11 12
13 14 15 16 17 18 19
20 21 22 23 24 25 26
27 28 29 30

Wk 11 • Day 072-294

March
Thursday

12

② lock (L) {

int acquired = 0;

while (!acquired) {

Cli();

if (L == 1) {

acquired = 1;

L = 0

}

sti();

}

}

unlock (L);

L = 1

}

⇒ Yes, this will work.

05

In this case, after thread 1 checks the condition at A and moves to step B, it's enables preemption, so it will not let any other thread unlock the lock.

Because of this cli() command, race condition does not occur.

08

Also if we recall, this locking discipline ensures race-freedom because following condition holds (as mentioned in lecture on locking)

- (1) We associate a lock with shared region (i.e. L)
- (2) Before accessing shared region, ensure corresponding associated lock is held.

This conditions are fulfilled here.

APRIL

MAY

JUNE

2020

13

March
Friday

Wk 11 • Day 073-293

	MARCH 2020					
Wk	M	T	W	T	F	S
09	30	31				1
10	2	3	4	5	6	7
11	9	10	11	12	13	14
12	16	17	18	19	20	21
13	23	24	25	26	27	28

• Assignment: Sleep and wakeup

09

Solⁿ: This implementation is correct.

It is okay if both producer (pcqwrite) and consumer (pcqread) are sleeping on the same channel q.

Even if the producer calls wakeup(q);

consumer does not have authority to

acquire the lock before producer released it.

And producer releases lock only after ~~wakeup~~ wakeup call.

Since this is Spinlock, consumer will be kept waiting unless producer released the lock.

04

So this implementation is good.

05

06

07

08



APRIL 2020						
Wk	M	T	W	T	F	S
14	1	2	3	4	5	
15	6	7	8	9	10	11
16	13	14	15	16	17	18
17	20	21	22	23	24	25
18	27	28	29	30		

2020

March
Saturday

14

Wk 11 • Day 074-292

Assignment: XV6 file system

- After adding the line

~~exit(0)~~)

printf ("log-write %d\n", b→sector/blockno);

① Output of echo a

⇒ log-write 34

log-write 34

log-write 59

② (I didn't get sector 29 as mentioned in homework.
Instead I got sector 59)

③ Output of echo X79

⇒ log-write 58

log-write 571

log-write 571

log-write 34

log-write 571

log-write 34

④ Output of rm a

⇒ log-write 59

log-write 34

log-write 58

log-write 34

log-write 34

Explanation for this are on next page.

APRIL

MAY

JUNE

2020

15

March
Sunday

Wk 11 • Day 075-291

09	30	31						
10	2	3	4	5	6	7	8	
11	9	10	11	12	13	14	15	
12	16	17	18	19	20	21	22	
13	23	24	25	26	27	28	29	

① Explanation for echo 79.

- In file creation, this steps happen
- ① Allocate an inode (write to inode region)
 - ② Update the size/status of the inode
(write to inode region)
 - ③ Add a directory-entry record to parent's directory's data blocks
(write to data block region)

∴ in log-write 34

log-write 34

log-write 59

34 is inode region

59 is data block region

② Exp. for echo X>9

The step happened are as follows -

1. Allocate a block (write to block bitmap region)
2. Zero out block contents (data block region)
3. Write "X" to block (to block data region)
4. update "a"s inode (size, addr)

∴ In this case; log-write 58

Indicates that 58 is block bitmap region

~~Block~~ "X" was written to data block region 57

two entries of log-write 34 = One for size update
One for addr update

APRIL 2020						
Wk	M	T	W	T	F	S
14		1	2	3	4	5
15	6	7	8	9	10	11
16	13	14	15	16	17	18
17	20	21	22	23	24	25
18	27	28	29	30		

2020

March
Monday

16

Wk 12 • Day 076-290

09 ① Explanation for rm a

10 → following steps occur:

11 ~~Allocate a block~~

- 1- Write zero to "a"s directory-entry record in parent directory's data blocks (write to data block region)
- 01 2- Update parent directory's inode (size)
- 02 3- Update "a"s inode to mark it free
- 02 4- update free block bitmap to mark data blocks free.

03 ∴ In output of rm a

04 In log-write 58, "a"s block bitmap region was freed.

05 In log-write 34, "a"s inode was marked free.

06 In log-write 59, zeroes were written to data block region.

07 ● Assignment ZCAV (plot attached for disk on physical laptop)

APRIL

MAY

JUNE

