

## Homework - 4

### Assignment: spinlocking

- i) The kernel panicked because ide lock is a spin lock. And spinlocks in xv6 should always run with interrupts disabled, if not in this case in our case panic will result because the interrupt handler would want to acquire the lock resulting in a deadlock.
- ii) The kernel didn't panic in this case of file-table lock because it is a sleep lock that doesn't need interrupts disabled and releases the lock before yielding. So a deadlock can't occur in this case. The difference in behaviour of file-table lock and ide lock is because of they are diff types of locks. ide lock is spinlock while file-table lock is sleep lock.
- iii) release() clears out lk->pu[0] and lk->qu[0] while holding the lock to preserve atomicity and prevent interleaving so that these instructions don't change the value from some other process while same other thread is holding the lock.

### Uniprocessor Locking

- i) This won't work on a uniprocessor because interrupts have been disabled before while sleep. This will result in the sleep running endlessly and no other process will get a chance to execute.
- ii) Yes, this will work on a uniprocessor as it will

give other processes a chance to execute by which can release or acquire the lock.

### Assignment: sleep and wakeup

Yes both the operations sleeping on the same channel is correct because when producer or consumer finishes its operation it would like to wake up the consumer or producer respectively who are sleeping on the same channel only. If they sleep on diff channels then would be an issue because on wakeup (as it wakes up by producer) it wakes up all the producers consumers sleeping on that channel by going through the process table and one of them will be able to execute.

Unrelated part of the code can call wakeup on a channel and then the consumer would check if its condition is satisfied if it still isn't then it would go to sleep again.

### Assignment: xv6 file system

echo > a

1 log-write 34  
log-write 34  
log-write 59

The first is allocating an inode on device driver marking it allocated and in the second is copying a modified in-memory inode to disk.  
The third is writing data to inode.

echo x > a.

old 1 log-write 58. from lmap start is 58 scattered to the first block  
2 log-write 666  
3 log-write 666  
4 log-write 34  
5 log-write 666  
6 log-write 34

1 → allocating a-zeroed block

2 → zeroing the entire block

3 → writing data to inode

4 → copy in memory inode to disk

5 → write data to inode

6 → copy in memory inode to disk

Two both exist because first old block is allocated then a is written

This can be seen as after first inode  $\rightarrow$  size is regular but after second it is 1.8

new a

1 log-write 59.  
2 log-write 34  
3 log-write 58.  
4 log-write 34  
5 log-write 34

1 - writing data to inode.

2 - copy in memory inode to disk

3 - free a disk block.

4 & 5  $\rightarrow$  modified in memory, inode to disk.

Assignment : ZCAV

Physical laptop disk : HP X360, i5 8<sup>th</sup> gen

TOSHIBA MQ04AB100 1 TB HDD

heads  $\geq 16$

Cycle time  $\rightarrow 120 \text{ ns}$

physical sector size  $\rightarrow 4096$

USB drive

Sandisk glide 16 GB

Physical laptop disk  $\rightarrow$  average throughput  $\approx 2$   $\rightarrow 97.5^-$   
92.8 MB in 95.11 seconds and 98.25 seconds.

Laptop disk

The zones in hard disks are around 20. In this HDD more than that around 30 are observed which could be due to less number of heads.

The max is around 140 mb/s (throughput)

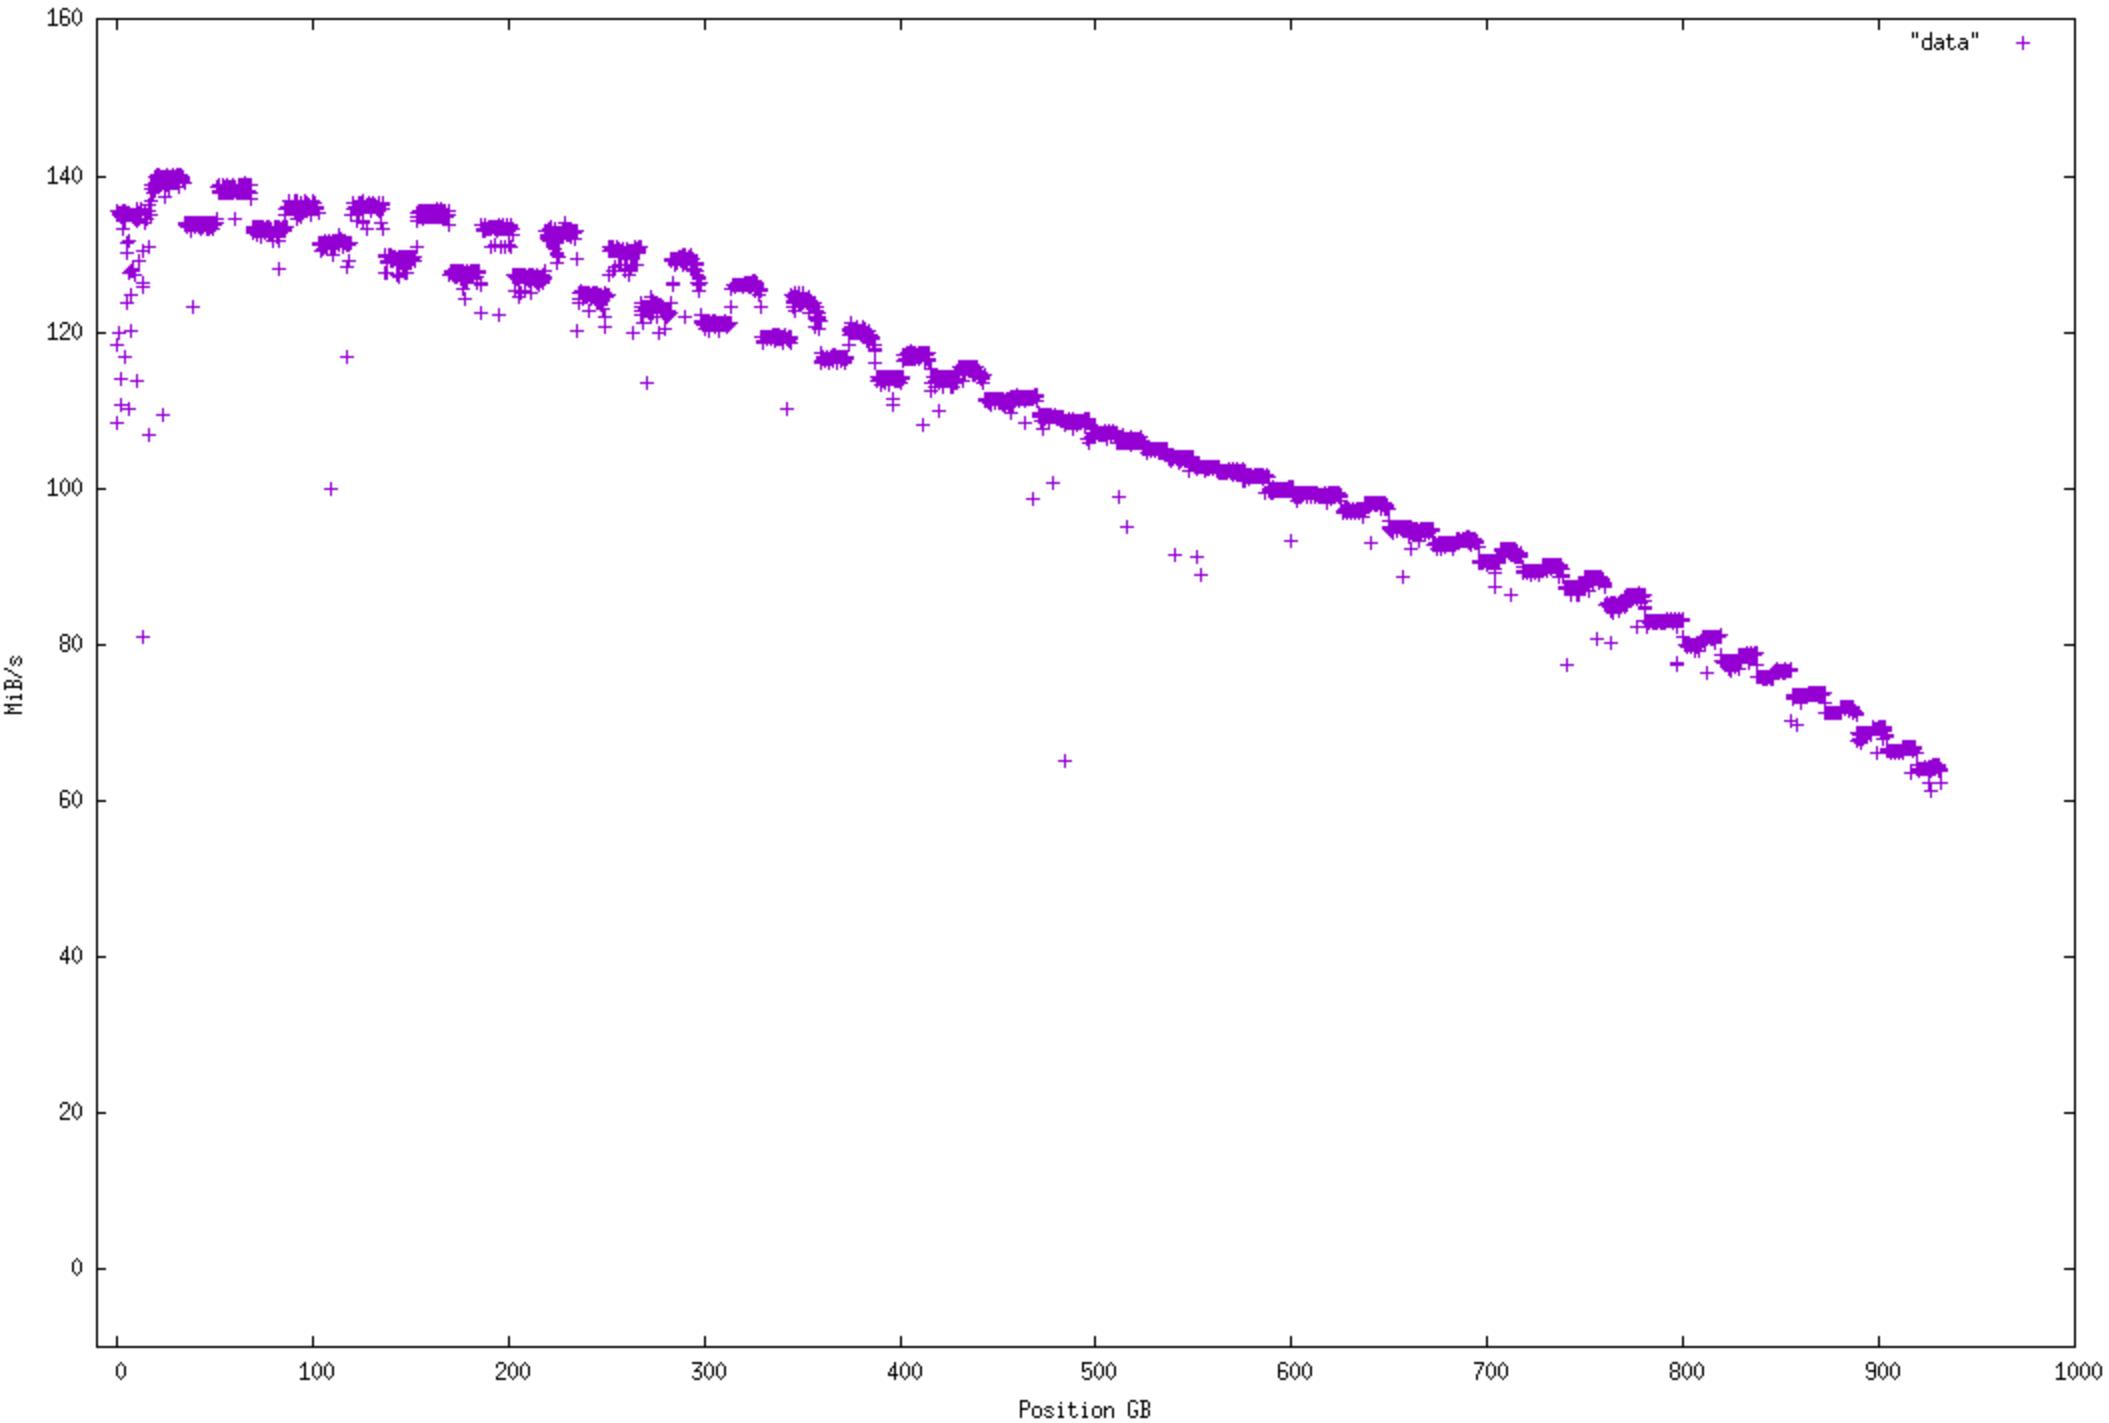
and min around 70 mb/s (throughput)

a drop of 50%.

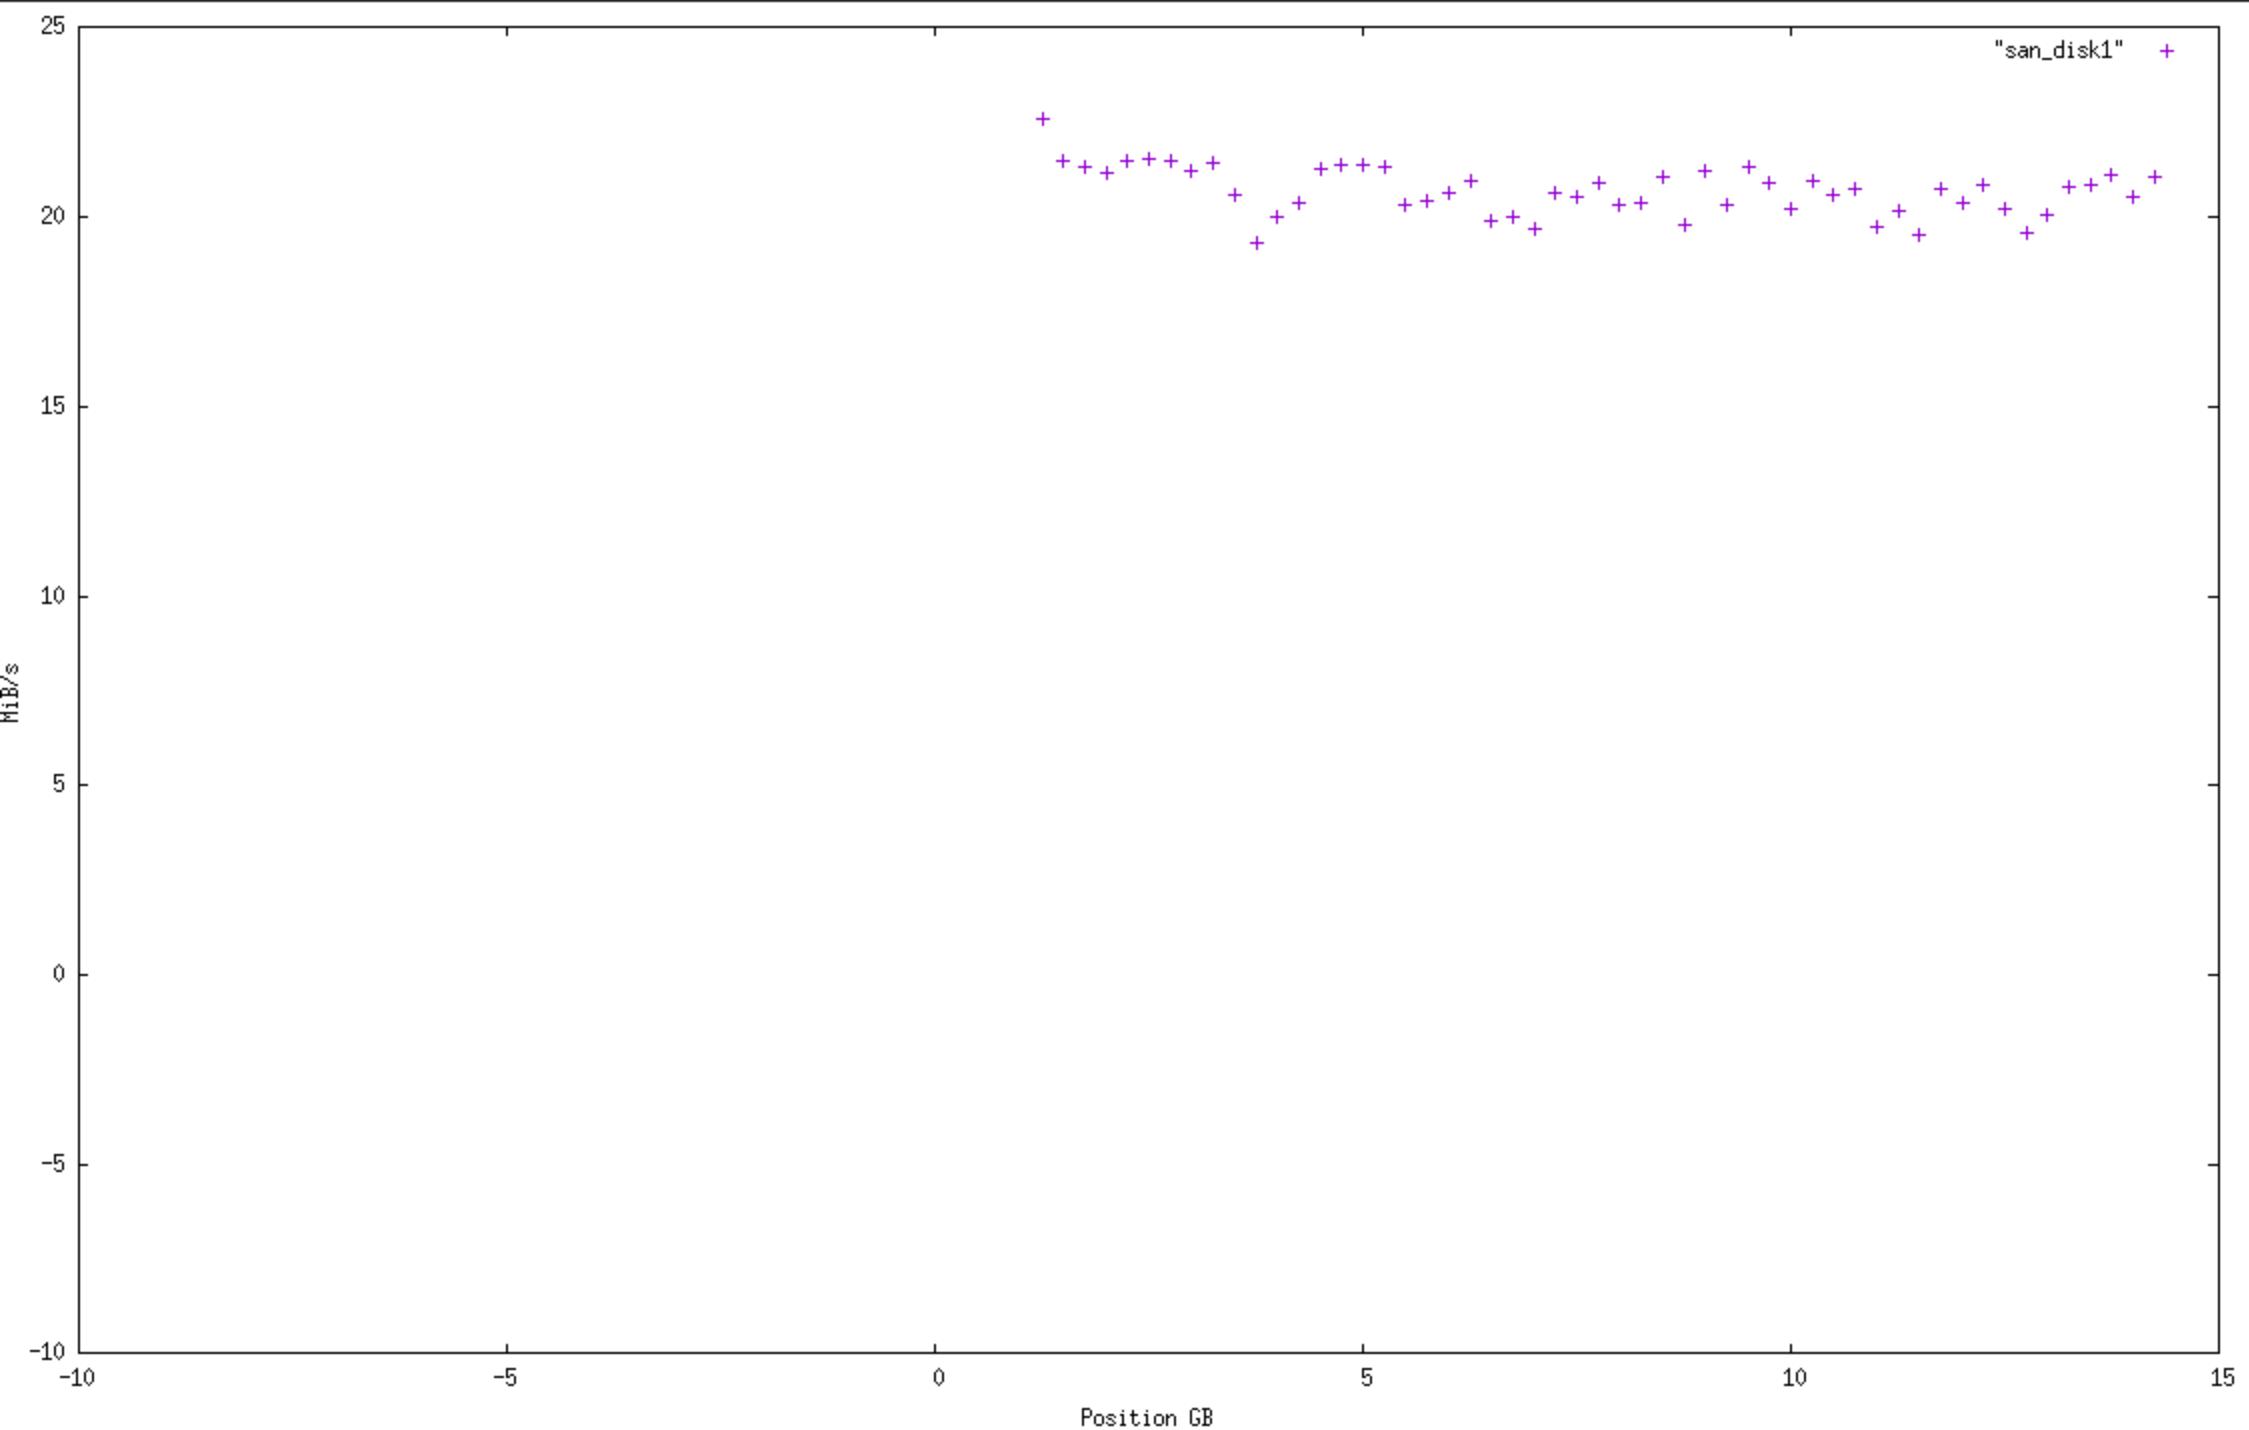
USB Disk

It has no moving parts and zones. So an almost constant throughput is observed of around 20 MB/s.

The plots are attached the one with label "data" is for HDD and the second one for USB.



608.615, 131.115



1,09173, 20,6378